

NLP on tweets: Analyzing twitter data about sexual abuse and assault

Prabha Nayak
December 3, 2019



1. Accessing Twitter public API to obtain data

- Create a twitter account if you do not already have one.
- Go to <https://apps.twitter.com/> and log in with your twitter credentials.
- Click "Create New App"
- Fill out the form, agree to the terms, and click "Create your Twitter application"
- In the next page, click on "API keys" tab, and copy your "API key" and "API secret".
- Scroll down and click "Create my access token" and copy your "Access token" and "Access token secret".

Number of tweets gathered = 19862

Number of features = 68

```

api_key <- 'otgCiwksUOZSwfzQEOikXWCZ4'
api_secret <- 'ib6QiWKZaCSKf1dHo1c1f34lGfFnCSys9COprmgjcrz087DrkQ'
access_token <- '1178855384288641024-DXojq1lXxBUTEMcBwHqNaNnrRY0F8o'
access_token_secret<- 'tkSiUU1jrw2GyaxiaAK9YFvqH6t23lZvJot8anYhHsBSk'

install.packages("twitter")
library(twitter)
#twittR does not include the capability of getting tweets >140 characters
install.packages("rtweet")
library(rtweet)

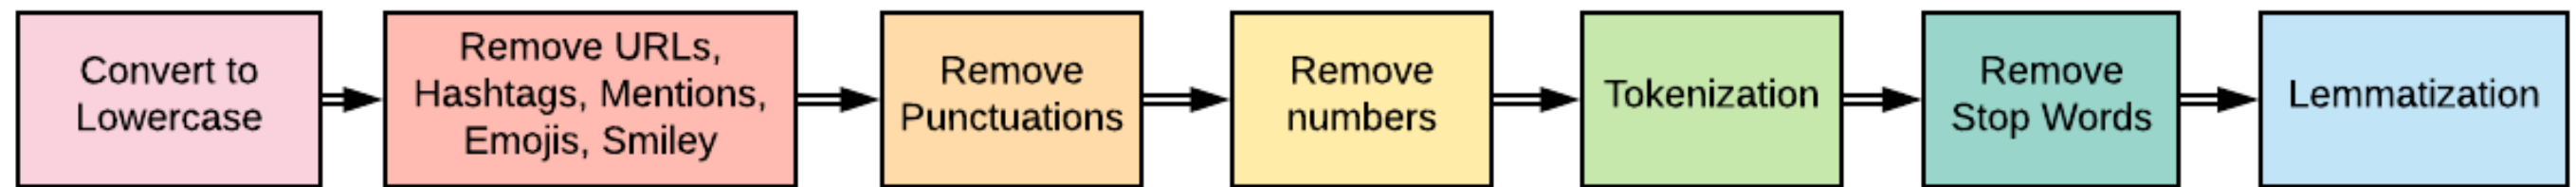
setup_twitter_oauth(api_key, api_secret, access_token, access_token_secret)

#Getting tweets
tweets <- search_tweets("sexual abuse -filter:retweets", n = 10000, lang = 'en', tweet_mode= 'extended')
tweets_assault <- search_tweets("sexual assault -filter:retweets",n=10000, lang='en', tweet_mode='extended')
#no_retweets <- strip_retweets(tweets)
tweetsdf <- tweets_data(tweets)
write_as_csv(x = tweetsdf, file_name = "C:/Users/PG/Desktop/NLP final/tweets_abuse.csv")
tweetsdf1 <- tweets_data(tweets_assault)
write_as_csv(x = tweetsdf1, file_name = "C:/Users/PG/Desktop/NLP final/tweets_assault.csv")

#Merging the 2 dataframes
all_tweets<- rbind(tweetsdf, tweetsdf1)
#tweetsdf <- tweetsdf[-c(2,5)]
write_as_csv(x = all_tweets, file_name = "C:/Users/PG/Desktop/NLP final/alltweets.csv")

```

NLP Pipeline



2. Preprocessing of data

a. Convert to lowercase

- Column 'text' extracted from the all_tweets dataframe
- Text converted to lowercase using lower()

```
23 all_tweets = pd.read_csv(r'C:\Users\PG\Desktop\NLP final\alltweets.csv')
24 print (all_tweets)
25
26 tweets = all_tweets[['text']]
27
28 disc_rows = tweets.drop_duplicates()
29
30 disc_list = disc_rows['text'].values.tolist()
31
32
33 lowertext = [x.lower() for x in disc_list]
34
35
```

2. Preprocessing of data

b. Remove URL, Punctuations and numbers

```
import preprocessor as p

newlist = []
punctuations = '!"#$%&'()*+,-./:;<=>?@[\\]^_`{|}~''''

for line in lowertext:
    line = p.clean(line)
    line = ''.join([i for i in line if not i in punctuations])
    line = ''.join([i for i in line if not i.isdigit()])
    newlist.append(line)
```

2. Preprocessing of data

c. Remove duplicates

- Each tweets ends with a distinct URL.
- So removing duplicates preferred after removing the URLs from tweets.

```
8 lowertext_disc = []  
9 [lowertext_disc.append(x) for x in newlist if x not in lowertext_disc]
```

2. Preprocessing of data

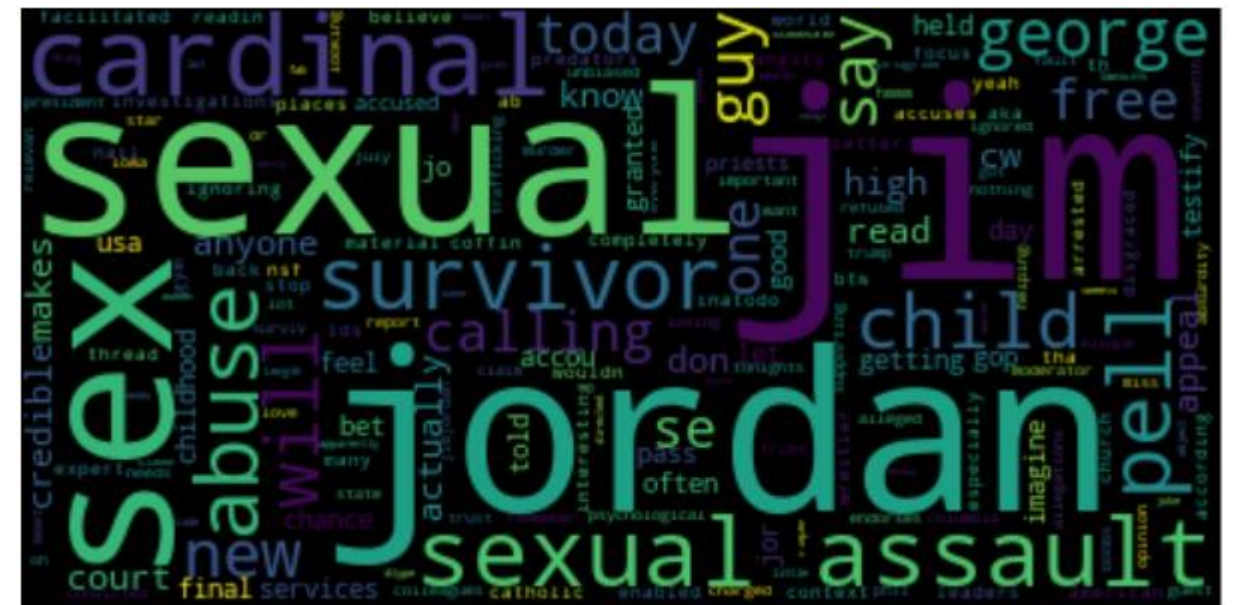
d. Tokenization

- Tokenization using the split() method

```
5 cleanText1 = pd.DataFrame(lowertext_disc, columns = ['tweets'])
6 cleanText1['tokens'] = ""

1 tokenized_text = cleanText1['tweets'].apply(lambda x: x.split())
2 cleanText2['tokens'] = tokenized_text
```


Word Cloud generation before



2. Preprocessing of data

e. Removing Stop Words

```
7 from nltk.corpus import stopwords
8 nltk.download('stopwords')

4 stop_words = set(stopwords.words('english'))
5
6 cleanText1['filtered-text'] = ""

9 no_stopWords = tokenized_text.apply(lambda x: [item for item in x if item not in stop_words])
```

2. Preprocessing of data

f. Lemmatization

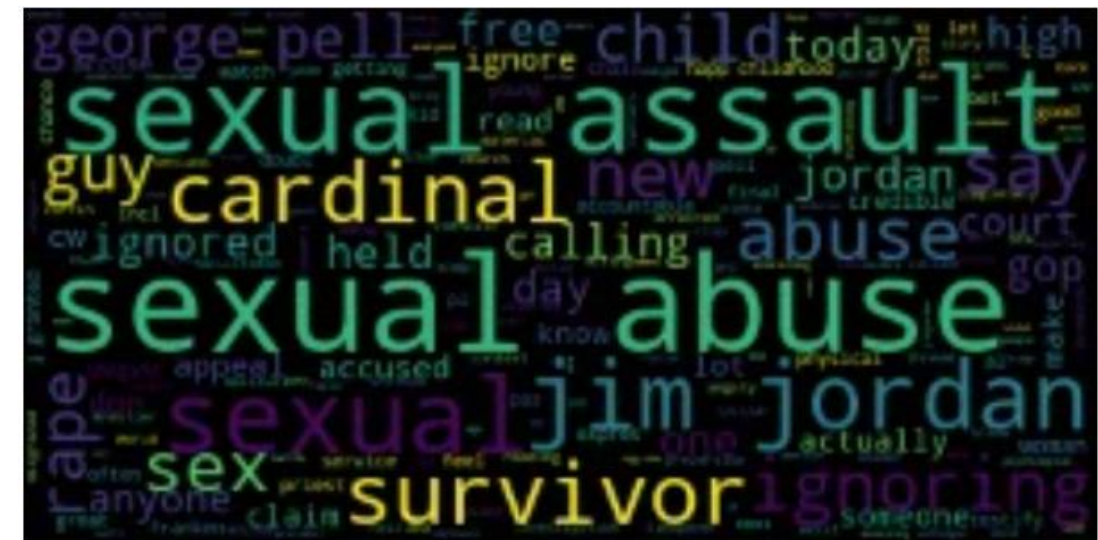
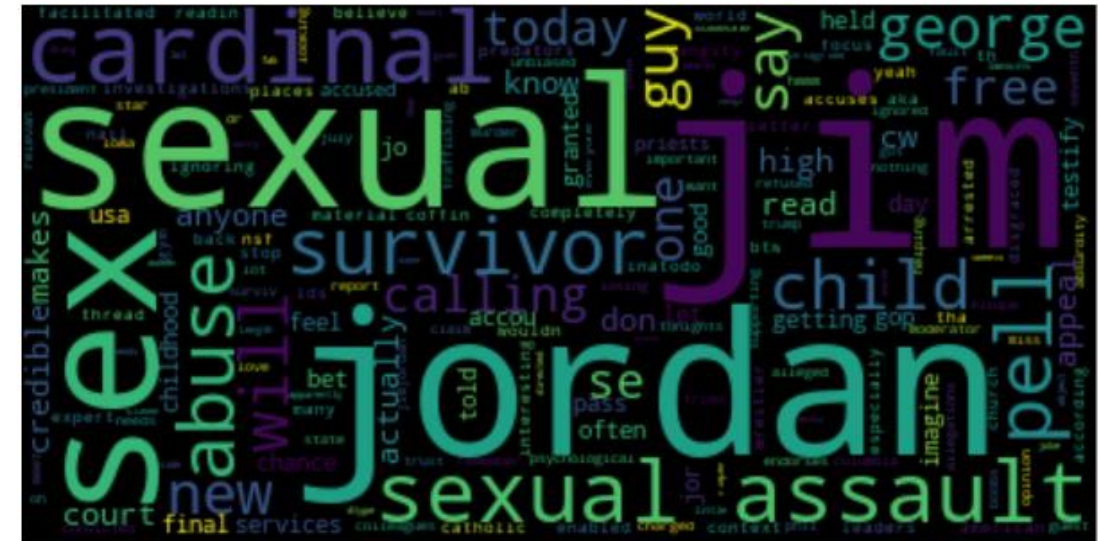
```
0 lemmatizer = WordNetLemmatizer()
1
2 def word_lemmatizer(text):
3     lem_text = [lemmatizer.lemmatize(i) for i in text]
4     return lem_text
5
6 lemmtext = cleanText1['filtered-text'].apply(lambda x:word_lemmatizer(x))
7
```

3. Analysis

a. Word Cloud generation

```
wordcloud = WordCloud().generate(str(tokenized_text))  
image = wordcloud.to_image()  
image.show()
```

```
wordcloud = WordCloud().generate(str(lemmtext))  
  
image = wordcloud.to_image()  
image.show()
```



3. Analysis

a. TF-IDF

```
8 def identity_tokenizer(text):
9     return text
0
1 tfidf = TfidfVectorizer(tokenizer=identity_tokenizer, lowercase=False)
2 vector1 = tfidf.fit_transform(no_stopWords)
3 features = tfidf.get_feature_names()
4
5
6 indices = np.argsort(tfidf.idf_)[::-1]
7
8 top_n = 50
9 top_features = [features[i] for i in indices[:top_n]]
0 print(top_features)
1
2 #frequency of the word sexual = 19823
3 tfidf.vocabulary_['sexual']
4
5 #frequency of the word assault = 1441
6 tfidf.vocabulary_['sexual']
7
8 #frequency of the word sexual = 96
9 tfidf.vocabulary_['abuse']
0
```

Next Steps

- Try n-grams
- Train-test split
- Fit training data to a model
- Fit the test data to the model to obtain results

References

1. <https://stackoverflow.com/questions/34860982/replace-the-punctuation-with-whitespace>
2. <https://codereview.stackexchange.com/questions/115954/writing-lists-to-csv-file>
3. <https://gis.stackexchange.com/questions/72458/exporting-list-of-values-into-csv-or-txt-file-using-arcpy>
4. <https://stackoverflow.com/questions/13464152/typeerror-unhashable-type-list-when-using-built-in-set-function/13464194>
5. <https://towardsdatascience.com/nlp-for-beginners-cleaning-preprocessing-text-data-ae8e306bef0f>
6. <https://stackoverflow.com/questions/48671270/use-sklearn-tfidfvectorizer-with-already-tokenized-inputs>
7. <https://medium.com/fintechexplained/nlp-text-processing-in-data-science-projects-f083009d78fc>
8. <https://www.geeksforgeeks.org/generating-word-cloud-python/>

Thank you!